

# Linear Regression

---

## 1) Code

```
def linear_function(theta_, x_) -> float:
    return np.inner(theta_, x_)

def linear_loss_all(theta_, X_, y_, N, lambd) -> float:
    sum = 0
    for i in range(N):
        new_x_ = np.append([1], X_[i])
        sum = sum + (linear_function(theta_, new_x_)-y_[i])**2
    return 0.5 * sum

...
The rest of the functions are the same for both regressions.
...

def find_graident(theta_, x_, y, function, lambd):
    new_x_ = np.append([1], x_)
    return (function(theta_, new_x_)-y) * new_x_ + (lambd * theta_)

def mini_batch_gradient_decent(
    theta_:np.ndarray,
    X_:np.ndarray,
    y_:np.ndarray,
    N:int,
    learning_rate:float,
    batch_size:int,
    function,
    lambd:float
):
    batch = random.sample(range(N), k=batch_size)
    gradient = sum([ find_graident(theta_, X_[i], y_[i], function, lambd)
                    for i in batch])
    new_theta_ = theta_ - ((learning_rate * gradient)/batch_size)
    return new_theta_

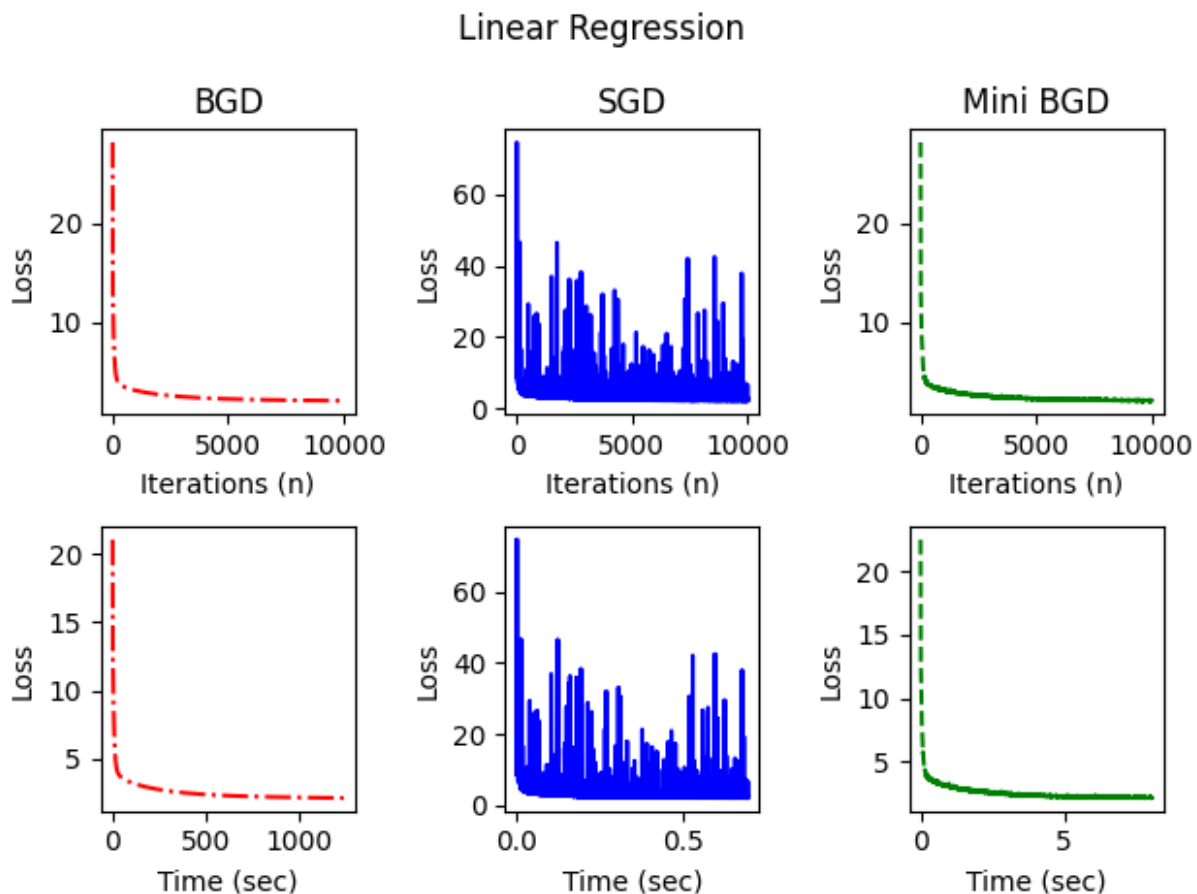
def gradient_decent(
    X_:np.ndarray,
    y_:np.ndarray,
    N:int,
    iterations:int,
    learning_rate:float,
    batch_size:int,
    function:str,
    lambd:float
):
```

```

function_types = {"linear":(linear_function, linear_loss_all), "logistic":(sigmoid_function, logistic_loss_all)}
f, j = function_types[function]
theta_ = np.array([0.001 for i in range(X_.shape[1]+1)])
losses = []
times = []
losses.append(j(theta_, X_, y_, N, lambd)/N)
for i in range(iterations):
    start = time.time()
    theta_ = mini_batch_gradient_decent(theta_, X_, y_, N, learning_rate, batch_size)
    end = time.time()
    losses.append(j(theta_, X_, y_, N, lambd)/N)
    times.append(end - start)
return (losses, times)

```

## 2) Plots



## 3) Findings relative to the convergence theory

I found that a bad guess for the initial weights/ $\theta$  would cause my regression to never converge. After some testing, I found that setting 0.001 for all  $\theta_i$ , the result of the regression would converge.

In addition to the initial weights, the learning rate/ $\alpha$  heavily impacted whether the regression would converge. Setting it to 0.00000005 helped the results converge the fastest without diverging.

# Logistic Regression

---

## 1) Code

```
def sigmoid_function(theta_, x_) -> float:
    return np.power(1 + math.exp(-1 * np.inner(theta_, x_)), -1)

def logistic_loss_all(theta_, X_, y_, N, lambd) -> float:
    sum = 0
    for i in range(N):
        reg = 0
        for theta in theta_:
            reg += theta**2
        new_x_ = np.append([1], X_[i])
        sum = sum - (
            y_[i] * math.log(sigmoid_function(theta_, new_x_)) +
            (1-y_[i]) * math.log(1-sigmoid_function(theta_, new_x_))
        ) + (lambd * reg)
    return sum

'''
The rest of the functions are the same for both regressions.
'''

def find_gradient(theta_, x_, y, function, lambd):
    new_x_ = np.append([1], x_)
    return (function(theta_, new_x_)-y) * new_x_ + (lambd * theta_)

def mini_batch_gradient_decent(
    theta_:np.ndarray,
    X_:np.ndarray,
    y_:np.ndarray,
    N:int,
    learning_rate:float,
    batch_size:int,
    function,
    lambd:float
):
    batch = random.sample(range(N), k=batch_size)
    gradient = sum([ find_gradient(theta_, X_[i], y_[i], function, lambd)
                     for i in batch])
    new_theta_ = theta_ - ((learning_rate * gradient)/batch_size)
    return new_theta_

def gradient_decent(
    X_:np.ndarray,
    y_:np.ndarray,
    N:int,
    iterations:int,
```

```

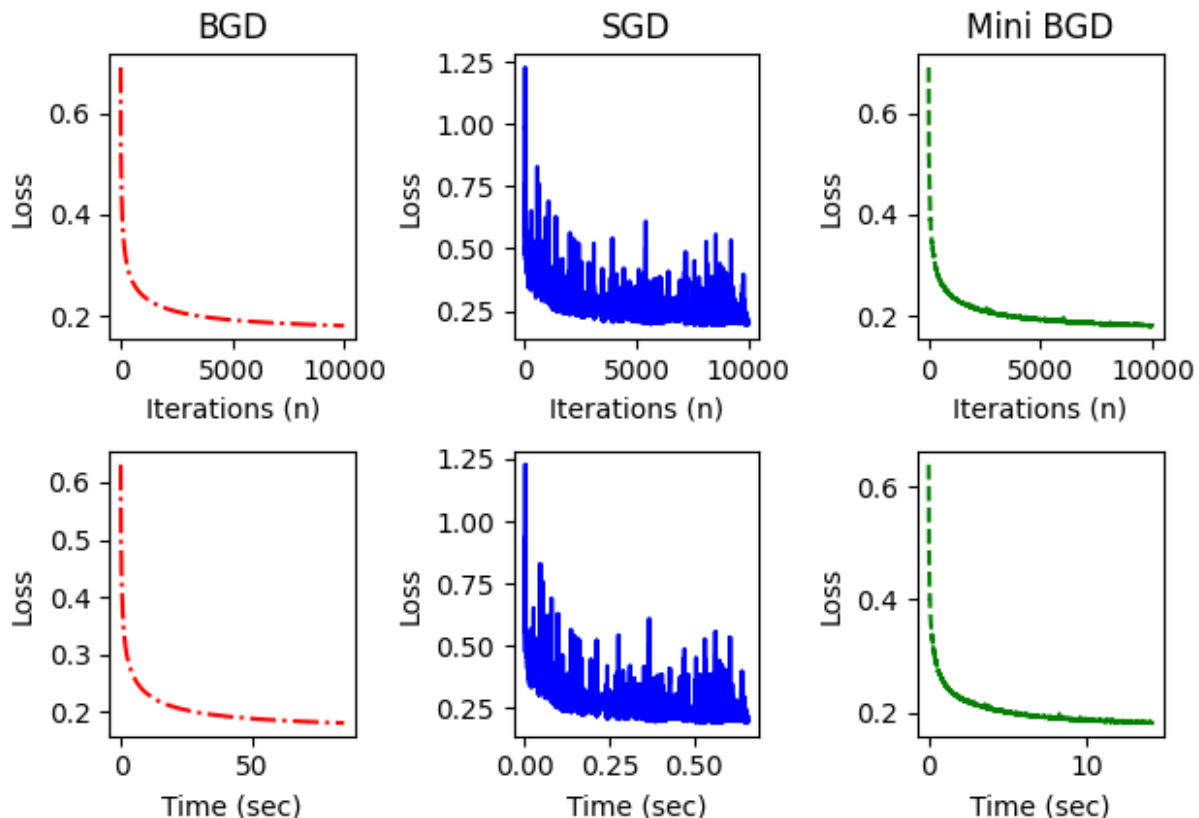
learning_rate:float,
batch_size:int,
function:str,
lambd:float
):
function_types = {"linear":(linear_function, linear_loss_all), "logistic":(sigmoid_function, logistic_loss_all)}
f, j = function_types[function]
theta_ = np.array([0.001 for i in range(X_.shape[1]+1)])
losses = []
times = []
losses.append(j(theta_, X_, y_, N, lambd)/N)
for i in range(iterations):
    start = time.time()
    theta_ = mini_batch_gradient_descent(theta_, X_, y_, N, learning_rate, batch_size)
    end = time.time()
    losses.append(j(theta_, X_, y_, N, lambd)/N)
    times.append(end - start)
return (losses, times)

```

## 2) Plots

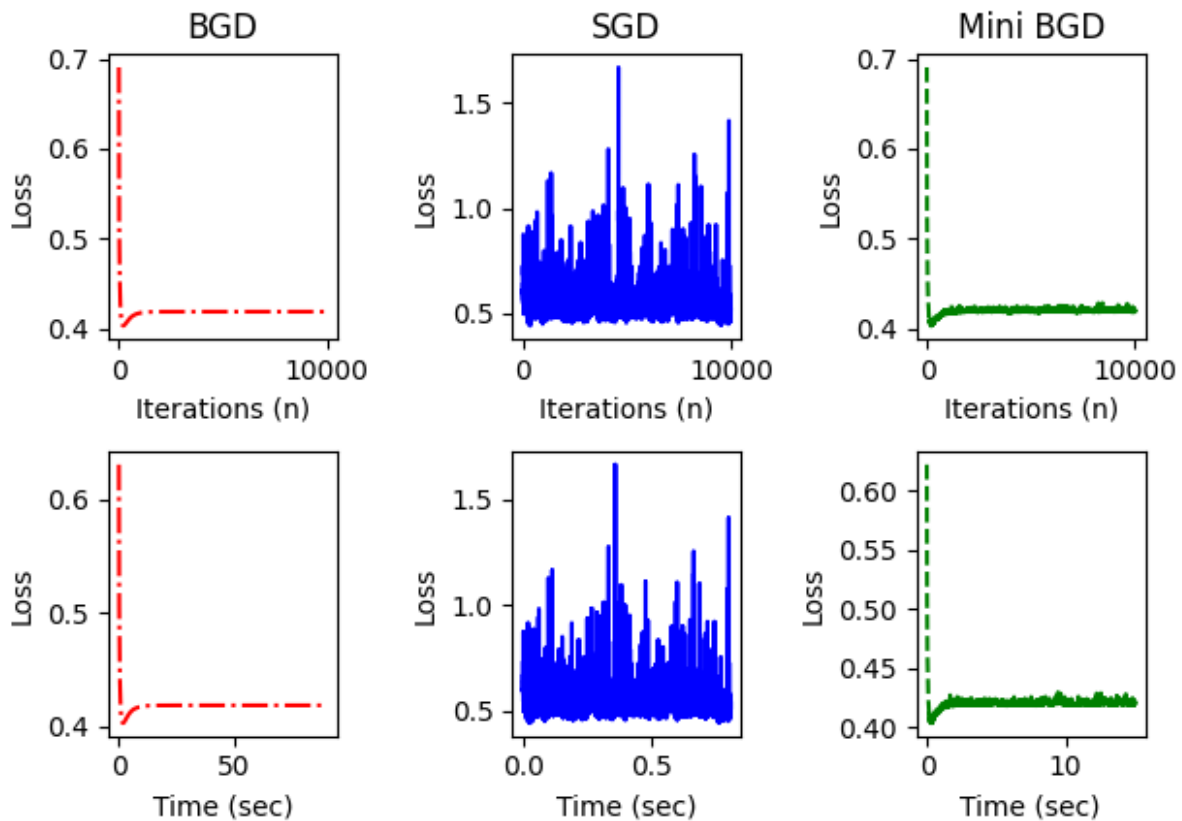
2a)

Logistic Regression,  $\lambda = 0$



2a)

## Logistic Regression, lambda = 0.01



### 3) Findings relative to the convergence theory

Like linear regression, I found that a bad guess for the initial weights/ $\theta$  would cause my regression to never converge. After some testing, I found that setting 0.001 for all  $\theta_i$ , the result of the regression would converge.

In addition to the initial weights, the learning rate/ $\alpha$  heavily impacted whether the regression would converge. The results from using 0.00000005 from linear regression showed that logistic regression converged slowly. Setting  $\alpha$  to 0.2, sped up the process and did not diverge.

### The rest of the code:

This contains imports, misc helper functions, formatting data, and displaying results of regressions.

```
import math
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import random
import time

def generate_ionosphere_data():
    sheet = pd.read_csv("ionosphere.csv")
    X_ = sheet.to_numpy()
    ..
```

```

y_ = X[:,34]
for i in range(y_.shape[0]):
    if y_[i] == "g":
        y_[i] = 1
    else:
        y_[i] = 0
X_ = np.delete(X_, np.s_[34:35], axis=1)
return (X_, y_)

def generate_air_quality_data():
    sheet = pd.read_excel("AirQualityUCI.xlsx")
    X_ = sheet.to_numpy()
    X_ = np.delete(X_, np.s_[0:2], axis=1)

    for i in range(X_.shape[1]):
        idx_to_change = []
        sum_ = 0
        for j in range(X_.shape[0]):
            val = X_[j][i]
            if (val == -200):
                idx_to_change.append(j)
            else:
                sum_ += val

        avg = sum_ / (X_.shape[0]-len(idx_to_change))
        for j in range(len(idx_to_change)):
            X_[idx_to_change[j]][i] = avg

    y_ = X[:,3]
    X_ = np.delete(X_, np.s_[3:4], axis=1)
    return (X_, y_)

def cumulate(times) -> None:
    for i in range(1, len(times), 1):
        times[i] = times[i] + times[i-1]

def run_regression(func:str):
    n = 10000
    itr = range(n+1)
    if (func == "linear"):
        learning_rate = 0.00000005
        X_, y_ = generate_air_quality_data()
    elif (func == "logistic"):
        learning_rate = 0.2
        X_, y_ = generate_ionosphere_data()
    else:
        exit(1)
    lambd = 0.01
    N = X_.shape[0]

    figure, axis = plt.subplots(2,3)

    losses, times = gradient_decent(X_, y_, N, n, learning_rate, N, func, lambd)

```

```

cumulate(times)
axis[0, 0].set_title("BGD")
axis[0, 0].plot(itr, losses, 'r-.', label='BGD')
axis[0, 0].set_xlabel="Iterations (n)", ylabel="Loss"
axis[1, 0].plot(times, losses[1:], 'r-.', label='BGD')
axis[1, 0].set_xlabel="Time (sec)", ylabel="Loss"

losses, times = gradient_decent(X_, y_, N, n, learning_rate, 1, func, lambda)
cumulate(times)
axis[0, 1].set_title("SGD")
axis[0, 1].plot(itr, losses, 'b-', label='SGD')

axis[0, 1].set_xlabel="Iterations (n)", ylabel="Loss"
axis[1, 1].plot(times, losses[1:], 'b-', label='SGD')
axis[1, 1].set_xlabel="Time (sec)", ylabel="Loss"

losses, times = gradient_decent(X_, y_, N, n, learning_rate, 50, func, lambda)
cumulate(times)
axis[0, 2].set_title("Mini BGD")
axis[0, 2].plot(itr, losses, 'g--', label='Mini BGD')
axis[0, 2].set_xlabel="Iterations (n)", ylabel="Loss"
axis[1, 2].plot(times, losses[1:], 'g--', label='Mini BGD')
axis[1, 2].set_xlabel="Time (sec)", ylabel="Loss"

figure.suptitle("Logistic Regression, lambda = 0.01")
figure.tight_layout()
plt.show()

def main():
    run_regression("linear")
    run_regression("logistic")

```

---

The entire program can be found on [Github](#)